

Mayuri Gulave

Data Scientist | AI Engineer

gulavemayuri1701@gmail.com | +91 9307302439 | Pune, Maharashtra

[Portfolio](#) | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

AI Engineer and Data Scientist with hands-on project experience building production-style ML, Deep Learning, NLP, and Generative AI systems. Developed and deployed multiple LLM-based applications independently — including RAG pipelines, AI chatbots, and NLP classifiers — using LangChain, ChromaDB, Llama 3, and OpenAI/Groq APIs. Strong statistical foundation from MSc Statistics combined with practical skills in Python, TensorFlow, and Scikit-learn. Eager to contribute to a team building real-world AI solutions across Finance, Healthcare, and NLP domains.

WORK EXPERIENCE

Data Scientist Intern | SMCS-PSI Data Analytics (OPC) Pvt. Ltd.

Jan 2026 – Jul 2026 · Pune

- Built and optimized ML and deep learning models — handling data preprocessing, feature engineering, and end-to-end model training and evaluation across business and healthcare datasets.
- Engineered LLM and RAG-based AI solutions for real-world business problems, achieving model accuracy improvements of ~15% through iterative hyperparameter tuning.
- Collaborated cross-functionally with a team of 4 engineers to streamline model workflows, reducing pipeline processing time by approximately 20%.

Data Science Intern | Rubix

May 2025 – Dec 2025

- Delivered data science work for a high-end capital firm — scoped requirements, engineered features, and built predictive models processing 50,000+ financial records.
- Designed a Credit Risk Prediction Model (Logistic Regression, Random Forest, XGBoost) achieving 89% ROC-AUC, enabling accurate loan default classification.
- Collaborated with cross-functional teams to translate client business challenges into ML solutions, delivering 3 models within a 6-month engagement.

EDUCATION

Pune University

Master's degree in Statistics · 2022–2024 · CGPA: 7.0

Bachelor's degree in Statistics · 2018–2021 · CGPA: 7.3

TECHNICAL SKILLS

Generative AI & LLMs

LangChain, ChromaDB, OpenAI API, Groq API, Llama 3, Ollama, Hugging Face Embeddings, Prompt Engineering, RAG Pipelines, Vector Databases (ChromaDB, FAISS), Streamlit

Programming & Tools

Python, SQL, Power BI, Tableau, Git, Jupyter Notebook, MS Excel

ML & Deep Learning Libraries

NumPy, Pandas, Scikit-learn, TensorFlow, Keras, Matplotlib, Seaborn, NLTK, SpaCy

Machine Learning

Regression, Classification, Clustering, XGBoost, Random Forest, Model Evaluation, Cross-Validation

Deep Learning

ANN, CNN, DNN, LSTM, RNN, Transfer Learning

NLP

Text Classification, Embeddings, Semantic Search, LLMs, NLTK, SpaCy

Deployment

Flask, REST APIs, Streamlit

Cloud & Big Data

AWS (Basic), Google Cloud (Basic), Apache Spark (Learning)

Statistics

Hypothesis Testing, Regression Analysis, ANOVA, Probability, Statistical Modeling

PROJECTS

Enterprise Knowledge Assistant (RAG Chatbot) · [GitHub](#)

Generative AI / NLP · Python, LangChain, ChromaDB, Llama 3, Ollama, Streamlit

- Built an enterprise-grade RAG chatbot using Llama 3 (via Ollama) and LangChain, enabling users to query internal knowledge bases with accurate, context-grounded answers.
- Implemented document ingestion pipeline into ChromaDB vector store; achieved sub-second semantic retrieval across large document corpora.
- Deployed via Streamlit with a clean chat UI; zero dependency on paid APIs by running Llama 3 fully locally.

AI Customer Support Ticket Classifier · [GitHub](#)

NLP / GenAI · Python, Llama 3, Prompt Engineering, Streamlit

- Developed an AI-powered ticket classification system using Llama 3 and structured prompt engineering to automatically categorize support tickets by type and priority.
- Designed prompt chains to extract intent, urgency, and category from free-text tickets — reducing manual triage effort significantly.
- Built and deployed interactive Streamlit interface for real-time ticket input and classification output.

AI Interview Coach · [GitHub](#)

NLP / GenAI · Python, Phi-3, Ollama, Streamlit, Resume Parsing

- Built an AI-powered interview preparation platform using Phi-3 (via Ollama) that generates role-specific questions and evaluates candidate answers.
- Implemented resume parsing to personalize question sets based on candidate background.
- Deployed fully locally using Ollama — no cloud API dependency, making it privacy-safe for resume data.

Credit Risk Prediction Model · [GitHub](#)

Finance / Banking · Python, Pandas, NumPy, Scikit-learn, Logistic Regression, Random Forest, XGBoost

- Built an ML model to classify customers as low-risk or high-risk for loan default, trained on 50,000+ records with 17 financial and demographic features.
- Achieved 89% ROC-AUC using XGBoost with feature engineering and SMOTE-based class balancing; surfaced top 5 predictive features for business stakeholders.

Chest X-Ray Image Classification · [GitHub](#)

Healthcare / Medical Imaging · Python, TensorFlow/Keras, CNN, OpenCV, NumPy, Matplotlib

- Designed a CNN architecture trained on 5,000+ chest X-ray images; applied augmentation and preprocessing achieving 91% validation accuracy.
- Evaluated model using precision, recall, and F1-score; addressed class imbalance via augmentation and weighted loss.

CERTIFICATIONS

Artificial Intelligence Expert — IABAC

Data Science Foundation — IABAC